

Unit 4

Parul[®] University NAAC A++

DIGITAL LEARNING CONTENT

Non-Monotonic Reasoning

Type of Reasoning:

1. Monotonic
2. Non-Monotonics

Definition:

- Monotonic means something that does not vary or change
- Non-Monotonic means something which can vary according to the situation or condition.

Parul[®] University NAAC A++

DIGITAL LEARNING CONTENT

Non-Monotonic Reasoning

- Non-monotonic reasoning allows revising conclusions with new information.
- Think of it as a toolbox where you might remove a tool if it becomes irrelevant.
- This makes reasoning more flexible and adaptable to changing situations.



Parul[®] University NAAC A++

DIGITAL LEARNING CONTENT

Non-Monotonic Reasoning

- Non-monotonic Reasoning is the process that changes its direction or values as the knowledge base increases.
- It is also known as NMR in Artificial Intelligence.
- Non-monotonic Reasoning will increase or decrease based on the condition.
- Since that Non-monotonic Reasoning depends on assumptions, It will change itself with improving knowledge or facts.
- Example:
Consider a bowl of water, If we put it on the stove and turn the flame on it will obviously boil hot and as we will turn off the flame it will cool down gradually

Monotonic Reasoning

Imagine a train always moving forward on a track. Monotonic reasoning is similar. It works with a fixed set of truths and builds upon them. Adding new information never contradicts existing conclusions. Here are its key features:

- **Always Expanding Knowledge Base:** As you learn more, your conclusions expand. You don't discard previous truths.
- **Focuses on True Facts:** Monotonic reasoning relies on established facts and proven theorems.
- **Examples:**
 - Mathematical proofs: Once proven, a theorem remains true regardless of new discoveries.
 - Inheriting traits in logic programming: If all birds can fly, and Tweety is a bird, then we can conclude Tweety can fly. Even if we learn sparrows can't fly, it

doesn't affect Tweety being a bird.

Non-Monotonic Reasoning

Non-monotonic reasoning, on the other hand, is like navigating a winding path. New information can overturn previous assumptions. It's crucial for real-world scenarios where knowledge is incomplete or uncertain. Here's what defines it:

- **Dynamic Conclusions:** New information can invalidate prior conclusions.
- **Deals with Incomplete Information:** It reasons with assumptions and defaults that might change with new evidence.
- **Examples:**
 - Everyday reasoning: We often assume birds can fly. But finding a penguin forces us to revise that assumption for penguins.
 - Default rules in AI: A rule might state "birds fly unless flightless." Discovering a penguin challenges this default.

Feature	Monotonic Reasoning	Non-Monotonic Reasoning
Adding Information	Expands knowledge base	May revise conclusions
Certainty of Facts	Relies on true facts	Deals with assumptions
Real-World Applicability	Limited	More suitable for real-world scenarios

Choosing the Right Approach

The choice between monotonic and non-monotonic reasoning depends on the situation:

- Use monotonic reasoning for problems with definite and unchanging truths (e.g., mathematical proofs).
- Non-monotonic reasoning is ideal for situations with incomplete or uncertain information (e.g., everyday decision-making).

In essence, both reasoning styles play a role in artificial intelligence and our own thinking. Monotonic reasoning provides a solid foundation, while non-monotonic reasoning allows us to adapt to the complexities of the real world.

Logics for Non-Monotonic Reasoning

Default Logic:

- Introduced by Raymond Reiter in 1980.
- Handles default assumptions.
- Makes assumptions (defaults) in the absence of specific information.
- Example: "Normally, birds fly" (default rule).

Mechanism:

- Defaults: $\frac{\text{Prerequisite: Justification}}{\text{Conclusion}}$
- Example: $\frac{\text{bird}(X): \text{flies}(X)}{\text{flies}(X)}$

Use Cases:

- Common-sense reasoning.
- Inferences with typical properties.

Autoepistemic Logic

- Developed by Robert C. Moore in the 1980s.
- Reasoning about an agent's own beliefs.
- Reasons about its own knowledge.
- Example: "I don't know it's raining, therefore it probably isn't raining."
- Revise the conclusion upon seeing new information: "It's cloudy outside."

Use Cases:

- Self-reflective reasoning.
- Modeling agent knowledge and beliefs.




Default Logic: A Closer Look

Imagine a world with "defaults" – general rules that typically hold true but have exceptions.

Default logic formalizes this reasoning style. Here's a breakdown:

- **Structure of a Default Rule:**
 - **Precondition (P):** A statement that activates the rule. If P is true, the rule might apply.
 - **Justification (J):** The conclusion drawn if the precondition is true, absent any defeater.
 - **Defeater (D):** A statement that cancels the justification, even if the precondition holds.
- **Non-monotonic Inference:** The power lies in how defaults interact. We can derive tentative conclusions (based on justifications) that might be withdrawn if defeaters emerge.

Here's an illustration:

- **Rule 1:** P: Tweety is a bird. J: Tweety flies. D: Tweety is a penguin.
- **Rule 2:** P: It's sunny. J: I will go for a walk. D: It's raining.
- **Example 1:** We know Tweety is a bird (P1 is true). Based on Rule 1, we can conclude Tweety flies (J1). However, if we discover Tweety is a penguin (D1 is true), the conclusion is withdrawn.
- **Example 2:** It's sunny (P2 is true), so we might conclude we'll go for a walk (J2). But if it starts raining (D2 is true), the plan changes.

Default Logic in Action:

- **Reasoning about Birds:** We can build a network of default rules about birds (most birds fly, penguins don't, ostriches can't, etc.). New information about a specific bird (e.g., Tweety being a penguin) can revise conclusions about its flying ability.

Autoepistemic Logic (AEL): Reasoning About What You Know

AEL goes beyond simple facts. It incorporates an agent's knowledge and self-awareness into the reasoning process. Here's the key concept:

- **Modal Operators:** AEL uses symbols like "Knows" (K) and "Believes" (B) to represent the epistemic state of the reasoner (the agent doing the reasoning).
- **Example:**
 - Our knowledge base (KB) states "All birds fly."
 - We don't know if Tweety is a bird.

AEL allows reasoning like:

* "If I knew Tweety is a bird ($K(\text{bird}(\text{Tweety}))$), then I would know it flies ($K(\text{flies}(\text{Tweety}))$). Since I don't know if Tweety is a bird ($\sim K(\text{bird}(\text{Tweety}))$), I cannot conclude it flies ($\sim K(\text{flies}(\text{Tweety}))$).

- **Reasoning with Incomplete Information:** AEL is adept at handling situations where knowledge is limited. It allows for expressing what we don't know, making reasoning more nuanced.

3. Theorist Preference Approach (TPA): Ranking Competing Theories

TPA takes a different approach. It focuses on:

- **Theories:** Sets of propositions that represent possible explanations for a situation.
- **Credos:** Measures of believability assigned to each theory.
- **Core Idea:** TPA ranks theories based on their credos. New information can alter these rankings, potentially favoring a different explanation.

- **Example:**

- Theory 1 (T1): Tweety is a bird (birds typically fly). $\text{Credo}(T1) = 0.8$ (fairly believable).
- Theory 2 (T2): Tweety is in a cage (caged birds can't fly). $\text{Credo}(T2) = 0.9$ (highly believable).

Initially, T1 might be favored (birds generally fly). However, if we learn Tweety is in a locked cage, the high credo of T2 makes it the preferred explanation, overriding T1.





Circumscription:

- Circumscription, introduced by John McCarthy, is a technique used in non-monotonic reasoning to capture the idea that things are typically as simple as possible unless there's evidence to the contrary. It essentially allows us to make inferences by minimizing the number of "special cases" that need to be considered.
- There are two main types of circumscription:
 - **Predicate Circumscription:** This focuses on a specific predicate (a property or relation) in your knowledge base. It assumes that an entity satisfies that predicate only if there's a reason for it to do so based on the existing knowledge. In simpler terms, it avoids making assumptions about an entity having a property unless there's explicit evidence for it.
 - **Domain Circumscription:** This type deals with the entire domain of discourse. It presumes that all the objects we're reasoning about are already accounted for in our knowledge base. There are no hidden entities waiting to be discovered.

Circumscription, introduced by John McCarthy, is a specific type of non-monotonic reasoning used in logic and artificial intelligence. It deals with the idea that things are typically as simple as possible, and only introduces complexity when necessary. Here's a breakdown of circumscription:

Core Principle: Preferring Simpler Explanations

Imagine a world where a light turns on whenever you flip a switch. Circumscription reasons that there likely aren't hidden gremlins controlling the light – the simplest explanation is the switch directly controls the light.

- **Minimizing Predicates:** Circumscription works by minimizing the number of "extensions" for a specific predicate in logic. An extension is essentially the set of things the predicate is true for.
- **Example:** Let $P(x)$ be a predicate meaning "x has property P." Circumscription aims to minimize the number of things that satisfy P, favoring simpler explanations.

Formalization: A Bit More Technical

Circumscription's formal definition goes beyond the intuitive explanation. It involves:

- **Minimal Models:** These are models (interpretations) of a logical theory that satisfy all its formulas and have the fewest number of elements satisfying a particular predicate.
- **Circumscribed Formula:** This formula expresses the set of all minimal models of the original theory.

Applications of Circumscription

- **The Frame Problem in AI:** This is the challenge of representing and reasoning about what doesn't change in a dynamic world. Circumscription can help by assuming things stay the same unless explicitly stated otherwise.
- **Diagnosis Systems:** In medical diagnosis, circumscription can be used to identify the simplest explanation for a set of symptoms, helping narrow down potential causes.

Limitations of Circumscription

- **Choosing the Right Predicate:** Circumscription requires careful selection of the predicate to minimize. Choosing the wrong one can lead to incorrect inferences.
- **Oversimplification:** In some cases, the simplest explanation might not be the most accurate one. Circumscription might miss important complexities.

There are two main types of circumscription:

- **Predicate Circumscription:** This focuses on a specific predicate (a property or relation) in your knowledge base. It assumes that an entity satisfies that predicate only if there's a reason for it to do so based on the existing knowledge. In simpler terms, it avoids making assumptions about an entity having a property unless there's explicit evidence for it.
- **Domain Circumscription:** This type deals with the entire domain of discourse. It presumes that all the objects we're reasoning about are already accounted for in our knowledge base. There are no hidden entities waiting to be discovered.

			
Comparison of Logics			
Logic Programming with Negation as Failure ➤ Used in Prolog. ➤ Assumes unprovable statements are false. Mechanism: Example: <code>flies(X) :- bird(X), not penguin(X)</code> Negation as Failure: $\neg p$ means "not provable p". Use Cases: ➤ Rule-based systems. ➤ Expert systems with default assumptions.	Key Differences: • Default Logic: Assumes typical properties. • Circumscription: Minimizes assumptions. • Autoepistemic Logic: Reflects on agent's beliefs. • Negation as Failure: Defaults to false if unprovable. • ASP (Answer set Programming): Uses stable models for problem-solving. • Abductive Logic: Generates best explanations. Use Case Suitability: • Varies based on problem domain and nature of information.		

Logic programming with negation as failure (often abbreviated NAF) is a powerful technique used in Artificial Intelligence (AI) for reasoning and problem-solving. It allows programmers to express rules with negation in a declarative way, even though classical logic doesn't directly handle negation as easily.

Here's a breakdown of the key concepts:

- **Logic Programming:** This is a programming paradigm based on logic. Programs consist of facts and rules, representing knowledge about the world.
- **Negation as Failure (NAF):** In NAF, the negation of a goal (statement) is considered true if attempting to prove the goal itself fails. It's not the same as classical "not" but a reasoning mechanism.

How NAF Works:

1. **Goal Statement:** The system is given a goal statement to prove.
2. **Rule Application:** It tries to find a rule in the program whose conclusion matches the goal.
3. **Subgoal Generation:** The rule's premises (conditions) become new subgoals.
4. **Recursive Process:** The system attempts to prove each subgoal using the same approach (steps 1-3).
5. **NAF and Failure:** If all attempts to prove a subgoal fail (no matching rule is found), then the negation of that subgoal is considered true (under NAF).

Example:

Imagine a program with these facts and rules:

- Fact: bird(tweety)
- Fact: can_fly(X) :- bird(X) (This rule says "X can fly if X is a bird")

Now, consider the goal: can_fly(tweety)

1. The system finds the rule matching the goal (can_fly(tweety)).
2. The rule's premise is bird(tweety), which is a fact, so it's proven.
3. Therefore, the original goal (can_fly(tweety)) is also proven.

Benefits of NAF:


- **Declarative Programming:** Allows programmers to focus on what holds true rather than how to achieve a specific outcome.
- **Elegant Negation Handling:** Provides a way to express negation without complex logical manipulations.
- **Reasoning Capabilities:** Enables programs to make logical deductions and draw conclusions.

Limitations of NAF:

- **Non-Monotonic Reasoning:** Conclusions based on NAF can change with new information. This can be counter-intuitive at times.
- **Efficiency Concerns:** Uncontrolled use of NAF can lead to infinite loops, requiring careful design and termination conditions.

Applications of Logic Programming with NAF:

- **Expert Systems:** Used to capture and reason with expert knowledge in a specific domain.
- **Natural Language Processing:** Can help understand the meaning of sentences and perform tasks like question answering.
- **Planning and Scheduling:** Aids in creating plans and schedules based on a set of goals and constraints.

Traditional Reasoning: Forward and Backward rules	Forward Chaining with Non-Monotonic Reasoning
➤ Non-monotonic reasoning can be integrated with both forward and backward chaining. ➤ Forward chaining: Needs to revisit conclusions when encountering new information. ➤ Starts with known facts, applies rules to reach conclusions (like dominoes falling). ➤ Backward chaining: Might involve making assumptions (defaults) to reach the goal. ➤ Works backward from a desired goal, searching for rules to justify it. 	➤ Starts with known facts and applies rules to reach conclusions (like dominoes falling). ➤ Needs to revisit conclusions when encountering new information that contradicts them. ➤ This is because new information might invalidate previously drawn conclusions. 

Forward Chaining Explained in Detail

Forward chaining, also known as data-driven reasoning, is a fundamental inference technique used in artificial intelligence (AI). It mimics a bottom-up approach to reasoning, starting with known facts and iteratively applying rules to reach new conclusions. Here's a deep dive into forward chaining with examples:

Core Principles:

- **Starts with Facts:** The process begins with a set of established facts or initial data points. These facts form the foundation upon which new knowledge is built.
- **Rule Application:** The system possesses a collection of rules. Each rule has a premise (IF clause) and a conclusion (THEN clause). Forward chaining iterates through these rules, checking if the premises of any rule match the existing facts.
- **Adding New Facts:** If a match is found, meaning all the conditions in the premise are true based on existing facts, the rule's conclusion is added as a new fact to the knowledge base.
- **Repetition:** This process of checking rules and adding new facts continues until:
 - No more rules can be applied because their premises aren't satisfied by the current set of facts.
 - A specific goal state is reached, indicating the desired conclusion has been derived.

Example 1: Diagnosing a Car Problem

Imagine a car diagnostic system using forward chaining:

- **Facts:** The car won't start (fact1) and the battery light is on (fact2).
 - **Rule 1:** IF the car won't start (premise1) AND the battery light is on (premise2), THEN the battery is dead (conclusion1).
 - **Rule 2:** IF the battery is dead (premise3), THEN jump-start the car (conclusion2).
1. The system starts with facts 1 and 2.
 2. It checks the rules. The premise of Rule 1 (fact1 and fact2) matches the existing facts.
 3. The conclusion of Rule 1 ("battery is dead") is added as a new fact.
 4. Since no new facts trigger further rules (premise3 of Rule 2 isn't met yet), the process might stop here. However, depending on the system's design, it could continue.

Example 2: Identifying a Customer's Eligibility for a Discount

Consider a customer discount system using forward chaining:

- **Facts:** The customer is a senior citizen (fact1) and they are purchasing groceries (fact2).
 - **Rule 1:** IF the customer is a senior citizen (premise1), THEN they are eligible for a senior discount (conclusion1).
 - **Rule 2:** IF the customer is purchasing groceries (premise2) AND they are eligible for a senior discount (premise3), THEN apply the senior discount (conclusion2).
1. The system starts with facts 1 and 2.
 2. It checks the rules. The premise of Rule 1 (fact1) matches a fact.
 3. The conclusion of Rule 1 ("eligible for senior discount") is added as a new fact.
 4. Now, the premise of Rule 2 (fact2 and the new fact) is also satisfied.
 5. The conclusion of Rule 2 ("apply senior discount") is reached.

Benefits of Forward Chaining:

- **Efficient for Event Detection:** It excels at continuously monitoring situations and identifying specific events based on incoming data.
- **Simple to Implement:** The logic flow is straightforward, making it easy to program

and understand.

- **Scalable:** It can handle large knowledge bases with many facts and rules efficiently.

Limitations of Forward Chaining:

- **Can Be Inefficient for Specific Goals:** If the desired goal is specific, the system might explore many irrelevant facts before reaching it. This can be computationally expensive.
- **Potential for Redundancy:** Processing all rules repeatedly might involve unnecessary checks if the goal is clear.

Real-World Applications:

- **Medical Diagnosis Systems:** Forward chaining can be used to analyze symptoms (facts) and apply medical knowledge (rules) to suggest potential diagnoses.
- **Network Intrusion Detection Systems (NIDS):** These systems monitor network traffic (data points) and use forward chaining rules to identify suspicious activity (conclusions).
- **Product Recommendation Systems:** Forward chaining can analyze customer purchase history (facts) and product attributes (rules) to recommend relevant products.

In conclusion, forward chaining offers a powerful approach for reasoning with data and dynamically building knowledge. By understanding its principles, benefits, and limitations, you can leverage this technique in various AI applications that require continuous learning and adaptation based on new information.

			
Backward Chaining with Non-Monotonic Reasoning ➤ Works backward from a desired goal, searching for rules to justify it. ➤ Might involve making assumptions (defaults) to reach the goal. ➤ These assumptions can be revised if new information contradicts them. 		Conclusion ➤ Forward chaining with revision: Incorporating non-monotonic reasoning, the system might need to revisit conclusions after applying a new rule. For instance, a forward-chaining system might conclude "Tweety flies" based on "Birds fly" and "Tweety is a bird." But upon encountering a rule "Penguins don't fly" and new information "Tweety is a penguin," it would need to revise its conclusion. ➤ Backward chaining with assumptions: Backward chaining with non-monotonic reasoning might involve making assumptions (defaults) to reach the goal. These assumptions can be revised if new information contradicts them. Imagine the system has a goal "Is the library open?" It might assume "Libraries are usually open on weekdays" unless it finds information about a specific closure.	

Backward Chaining Explained in Detail

Backward chaining, also known as goal-driven reasoning, is another fundamental inference technique used in artificial intelligence (AI). It takes a top-down approach, starting with a specific goal and working backward to identify the facts or conditions that must be true to achieve that goal. Let's delve deeper into backward chaining with examples:

Core Principles:

- **Starts with a Goal:** The process begins with a well-defined goal or hypothesis that you want to prove or disprove.
- **Rule Matching:** The system searches its knowledge base for rules where the conclusion (THEN clause) matches the desired goal.
- **Subgoal Generation:** If a matching rule is found, the premises (IF clause) of that rule become new subgoals. These subgoals represent smaller conditions that need to be satisfied for the rule's conclusion (the original goal) to hold true.
- **Recursive Process:** The system attempts to prove each subgoal by repeating steps 2-4. It keeps working backward, searching for rules that support the subgoals until:
 - All subgoals are successfully proven true, meaning the original goal is confirmed.
 - No matching rule is found for a subgoal, indicating the overall goal cannot be achieved based on the available knowledge.

Example 1: Crime Scene Investigation

Imagine a detective using backward chaining to investigate a robbery:

- **Goal:** Did John steal the money?
- **Rule 1:** IF John stole the money (conclusion1), THEN there should be fingerprints at the crime scene matching John's (premise1).
- **Rule 2:** IF there are fingerprints at the crime scene matching John's (conclusion2), THEN John was at the crime scene (premise2).

1. The investigation starts with the goal: Did John steal the money?

2. The system searches for a rule with "John stole the money" as the conclusion. Rule 1 matches.

3. The premise of Rule 1 ("fingerprints at the crime scene matching John's") becomes the first subgoal.
4. If fingerprint evidence confirms the subgoal (John's fingerprints are found), the system moves on.
5. The next subgoal would be the premise of Rule 2 ("John was at the crime scene"). The investigation would focus on finding evidence to support this.

Example 2: Trip Planning

Consider a travel planner using backward chaining to determine if a trip to Hawaii is feasible:

- **Goal:** Can I travel to Hawaii next month?
 - **Rule 1:** IF you can travel to Hawaii (conclusion1), THEN you have a valid passport (premise1).
 - **Rule 2:** IF you have a valid passport (conclusion2), THEN your passport doesn't expire before your trip (premise2).
1. The goal is to determine if a trip to Hawaii is possible.
 2. The system searches for a rule with "travel to Hawaii" as the conclusion. Rule 1 matches.
 3. The premise of Rule 1 ("valid passport") becomes the first subgoal.
 4. If you have a valid passport, the system checks the next subgoal based on Rule 2 (passport expiry).

Benefits of Backward Chaining:

- **Efficient for Specific Goals:** It's well-suited for situations where you have a clear goal in mind. The system focuses on what's necessary to achieve that goal, potentially avoiding unnecessary processing of irrelevant information.
- **Good for Diagnosis and Troubleshooting:** It's helpful in scenarios where you want to identify the root cause of a problem or confirm a hypothesis.
- **Can Handle Incomplete Information:** Backward chaining can work even if some subgoals cannot be proven definitively. It might suggest alternative paths or highlight uncertainties.

Limitations of Backward Chaining:

- **Can Be Inefficient for General Exploration:** If the goal is broad or poorly defined, the system might explore many rules without a clear direction, potentially leading to inefficiency.
- **Requires Well-Defined Rules:** The success of backward chaining hinges on having a comprehensive and well-organized knowledge base with clear rules and relationships.

Real-World Applications:

- **Expert Systems:** Backward chaining is used in expert systems to diagnose diseases, recommend financial products, or troubleshoot technical problems.
- **Game Playing AI:** AI players in games might use backward chaining to plan their moves by starting with a desired outcome (winning) and working backward to identify the necessary actions.
- **Theorem Proving:** In automated theorem proving, backward chaining can be used to establish the truth of a theorem by proving the implications of its negation.

In conclusion, backward chaining offers a powerful approach for goal-oriented reasoning and problem-solving. By understanding its principles, benefits, and limitations, you can leverage this technique in various AI applications that require focused reasoning and efficient exploration of possibilities to achieve specific goals.



Justification-Based Truth Maintenance Systems

Justification-Based Truth Maintenance Systems (JTMS) are a type of system used in Artificial Intelligence (AI) to manage the consistency of beliefs within a knowledge base. Unlike a traditional database that simply stores true or false statements, a JTMS focuses on the reasons why a belief is held. This allows the system to handle situations where new information might contradict existing beliefs.

The Challenge of Belief Management in AI :

- Traditional databases store information as true or false statements.
- AI systems need to deal with uncertainty and potentially inconsistent information.
- New information might contradict existing beliefs.

Justification-Based Truth Maintenance Systems (JTMS) Explained in Depth

Justification-Based Truth Maintenance Systems (JTMS) are a cornerstone of non-monotonic reasoning in Artificial Intelligence (AI). They offer a powerful way to handle situations where

beliefs can change as new information comes to light. Let's delve deeper into the intricate workings of JTMS:

Core Principles:

- **Beliefs and Labels:** JTMS operates on a network of beliefs. Each belief is represented by a node in the network. These nodes carry labels, typically "IN" (believed) or "OUT" (not believed). Imagine a detective's board with suspects and evidence – JTMS works in a similar fashion, tracking beliefs and their relationships.
- **Justifications:** Crucially, each belief is linked to a justification. A justification is essentially a rule or a combination of beliefs that explains why the main belief is considered true. Justifications are often depicted as Directed Acyclic Graphs (DAGs), where nodes represent beliefs and arrows depict dependencies between them. Think of justifications as the detective's reasoning behind connecting a suspect to the crime scene.
- **Maintaining Consistency:** A core function of JTMS is to ensure consistency within the belief network. This involves constantly monitoring justifications. If the justification for a belief becomes inconsistent (meaning a premise in the justification becomes labeled "OUT"), the system might re-evaluate the belief itself, potentially changing its label to "OUT" (not believed). Imagine the detective discovering an alibi for a suspect, forcing them to reconsider their initial belief about the suspect's involvement.

Parul [®] University NAAC A++		DIGITAL LEARNING CONTENT	
How JTMS works ➤ Belief Representation: Each piece of information is a node (circle). ➤ Labels: Nodes are labelled "IN" (believed) or "OUT" (not believed). ➤ Justifications: Rules connect beliefs (lines with arrows). If conditions are met (beliefs "IN"), the conclusion is "IN". ➤ Maintaining Consistency: The system revises labels when new information or retractions arise.		Advantages of JTMS ➤ Non-Monotonic Reasoning: Reasoning adapts based on new information. Conclusions aren't absolute and can be revised. ➤ Efficient Belief Revision: The system focuses on revising relevant beliefs based on justification dependencies. ➤ Reasoning about Assumptions: Analyze hypothetical scenarios by making assumptions and tracking consequences within the JTMS framework.	

Operations in a JTMS:

1. **Adding Beliefs:** When a new belief is introduced, the system attempts to find a justification that supports it. This justification could be a pre-defined rule or a combination of existing beliefs.
 - If a valid justification is found (all its premises are labeled "IN"), the new belief is added to the network and labeled "IN" (believed). This is like the detective adding a new suspect to the board based on a witness statement.
2. **Maintaining Consistency:** The system constantly checks for inconsistencies. If a belief's justification becomes invalid (a premise becomes "OUT"), the system reevaluates the belief. Depending on the JTMS implementation, the belief might be:
 - Labeled "OUT" (not believed anymore): This reflects the detective discovering new evidence that contradicts the initial belief about a suspect.
 - Quarantined for further investigation: This allows for exploring alternative justifications or gathering more information before discarding the belief entirely.
3. **Retracting Beliefs:** Sometimes, a belief needs to be explicitly retracted (no longer considered true). When this happens, the system investigates the justifications of other beliefs that depend on the retracted belief. Beliefs with justifications that are no longer valid due to the retraction might also be retracted or quarantined. Imagine the detective retracting their initial suspicion about a suspect, leading to a reevaluation of other beliefs based on that suspicion.

Benefits of JTMS:

- **Non-Monotonic Reasoning:** The key strength – JTMS allows for revising conclusions as new information becomes available. This is crucial in real-world scenarios where situations can change dynamically.
- **Explanation Facility:** JTMS can explain why a belief is held. By tracing the justification network, you can understand the reasoning behind a particular conclusion. This is like the detective presenting the evidence and reasoning that led them to suspect someone.
- **Efficiency:** JTMS focuses on revising only beliefs affected by changes in justification.

This avoids unnecessary re-evaluation of the entire belief network, making it computationally efficient for large systems.

Limitations of JTMS:

- **Complexity of Justifications:** Constructing and maintaining complex justifications, especially for intricate relationships between beliefs, can be computationally expensive.
- **Limited Explanations:** Explanations are limited to the justifications within the system itself. They might not capture the full reasoning process, especially if there are implicit assumptions or incomplete information.

Real-World Applications of JTMS:

- **Diagnosis Systems:** Medical diagnosis systems can utilize JTMS to track potential diagnoses and their justifications based on symptoms. New test results can then be incorporated, potentially revising diagnoses.
- **Planning and Scheduling Systems:** Planning systems can leverage JTMS to maintain consistent plans as new information about resource availability or task dependencies emerges.
- **Natural Language Processing (NLP):** NLP systems can use JTMS to reason about the meaning of sentences and revise interpretations as the context changes during a conversation.




DIGITAL LEARNING CONTENT



Semantic Nets

➤ **Function:** Capture and structure knowledge by portraying it as a network of interconnected nodes and edges.

Components:

- **Nodes:** Represent entities or concepts (like dog, cat, animal).
- **Edges:** Describe the relationships between those entities (like "is-a" - dog is-a animal).
- **Benefits:**
 - Intuitive representation: Easy to visualize knowledge structures.
 - Efficient reasoning: Enables reasoning tasks like inheritance (dogs inherit traits from animals)



Semantics Nets: Unveiling the Meaning Web

Semantic nets, also known as semantic networks, are a knowledge representation technique used in artificial intelligence (AI) and natural language processing (NLP) to model

relationships between concepts. Imagine a giant web of ideas interconnected by threads of meaning – that's the essence of a semantic net.

Core Principles:

- **Nodes and Arcs:** A semantic net is a graph-based structure. The building blocks are nodes, which represent concepts, entities, or objects in the world (e.g., "dog," "bark," "忠诚" (zhōng chéng, meaning "loyalty" in Chinese)). These nodes are connected by arcs, which represent relationships between the concepts (e.g., "has-sound" arc connecting "dog" and "bark").
- **Labeled Arcs:** The power of semantic nets lies in the labels on the arcs. These labels specify the nature of the relationship between the connected nodes. They can be simple labels like "is-a" (e.g., "dog is-a mammal") or more complex ones like "located-at" or "causes."
- **Inheritance:** Semantic nets often leverage inheritance, where properties of a general concept are inherited by more specific ones. For example, in a net with a node "animal" connected to a node "can-move" by a "has-ability" arc, all subclasses of "animal" (like "dog") would inherit the ability to move.

Example: A Simple Animal Semantic Net

Imagine a net with nodes for "animal," "dog," "cat," "bark," and "meow." The relationships might be represented by arcs labeled:

- "is-a" (dog is-a animal, cat is-a animal)
- "has-sound" (dog -has-sound-> bark, cat -has-sound-> meow)

Benefits of Semantic Nets:

- **Intuitive Representation:** They offer a visually appealing and easy-to-understand way to represent knowledge. The connections between concepts are clear and readily interpretable.
- **Efficient Reasoning:** Semantic nets enable reasoning by traversing the connections. You can infer new knowledge by following the arcs and their labels. For example, if "dog is-a mammal" and "mammal can-nurse," then you can infer that "dog can-nurse."
- **Scalability:** They can be easily extended by adding new nodes and arcs, making them

suitable for representing large and evolving knowledge bases.

Limitations of Semantic Nets:

- **Limited Expressiveness:** While powerful for basic relationships, they can struggle to capture complex relationships or nuances of natural language.
- **Disambiguation Challenges:** With homonyms (words with multiple meanings), it can be unclear which meaning is intended by a node without additional context.
- **Reasoning Complexity:** Reasoning tasks can become computationally expensive for very large and intricate semantic nets.

Applications of Semantic Nets:

- **Natural Language Processing (NLP):** Semantic nets are used in NLP tasks like word sense disambiguation (identifying the intended meaning of a word) and machine translation (understanding the meaning of text in one language and expressing it in another).
- **Information Retrieval:** Search engines can leverage semantic nets to understand the relationships between search terms and improve the relevance of search results.
- **Knowledge Representation:** They are used in various AI systems to represent knowledge about a specific domain, such as medical diagnosis systems or expert systems.



Statistical Reasoning in AI

- **Function:** Reason about and draw conclusions from data by leveraging statistical methods.
- **Methods:**
 - **Probability Theory:** Calculates the likelihood of events.
 - **Bayesian Networks:** Represent relationships between variables and their probabilities.
- **Benefits:**
 - Handles uncertainty: Deals with situations where information is incomplete or imprecise.
 - Learning from data: Can learn and improve reasoning based on new data.



Statistical reasoning is a fundamental approach in artificial intelligence (AI) that utilizes statistical methods to analyze data, identify patterns, and make predictions under uncertainty.

It plays a crucial role in enabling AI systems to learn from data, adapt to new information, and make informed decisions in real-world scenarios.

Here's a deeper dive into statistical reasoning in AI:

Core Principles:

- **Probability and Statistics:** Statistical reasoning heavily relies on the concepts of probability and statistics. Probability theory provides a mathematical framework to quantify the likelihood of events, while statistical methods offer tools for analyzing data and drawing inferences from it.
- **Data-Driven Approach:** Unlike traditional rule-based AI, statistical reasoning emphasizes learning from data. AI systems are trained on large datasets, allowing them to identify patterns and relationships within the data.
- **Uncertainty Handling:** Real-world data often contains inherent uncertainty. Statistical reasoning allows AI systems to account for this uncertainty and express their predictions as probabilities or confidence intervals. This provides a more nuanced understanding of the system's outputs.

Key Techniques:

- **Machine Learning:** A vast array of machine learning algorithms leverage statistical reasoning. These algorithms learn from data to perform tasks like classification (predicting a category for new data points), regression (predicting continuous values), and clustering (grouping similar data points).
 - **Example:** A spam filter uses statistical reasoning to analyze email content and identify patterns that distinguish spam from legitimate emails.
- **Probabilistic Inference:** Techniques like Bayesian networks and Markov models allow AI systems to represent relationships between variables and reason under uncertainty. This is useful for tasks like medical diagnosis or anomaly detection.
 - **Example:** A recommendation system might use probabilistic inference to recommend products to users based on their past purchase history and the buying habits of similar users.

Benefits of Statistical Reasoning in AI:

- **Data-Driven Learning:** Enables AI systems to learn and improve continuously as they

are exposed to more data.

- **Robustness and Adaptability:** Allows AI systems to handle new and unseen situations by generalizing from the patterns learned from data.
- **Uncertainty Quantification:** Provides a more realistic view of the system's performance by expressing predictions as probabilities.
- **Wide Applicability:** Statistical reasoning can be applied to a vast range of AI tasks, from computer vision and natural language processing to robotics and finance.

Limitations of Statistical Reasoning in AI:

- **Data Dependence:** The performance of statistical methods heavily relies on the quality and quantity of data. Biased or limited data can lead to biased or inaccurate results.
- **Interpretability Challenges:** Some machine learning models can be complex and their reasoning process might not be readily interpretable. This can be a challenge in situations where understanding the "why" behind a decision is crucial.
- **Computational Cost:** Training complex statistical models can be computationally expensive, especially for large datasets.

Real-World Applications of Statistical Reasoning in AI:

- **Self-Driving Cars:** Utilize statistical reasoning to analyze sensor data and make decisions about navigation, obstacle avoidance, and traffic light recognition.
- **Fraud Detection:** Financial institutions use AI to analyze transaction patterns and identify potential fraudulent activities based on statistical anomalies.
- **Medical Diagnosis:** AI systems can analyze medical images and patient data to assist doctors in diagnosing diseases with higher accuracy.




Combining Semantic Nets and Statistical Reasoning

- Merging the strengths of both approaches.
- Adding weights to edges in semantic nets: Represent the strength or likelihood of a relationship.
- Statistical reasoning to analyze data associated with concepts in a semantic net



Feature	Semantic Nets	Statistical Reasoning	Combining Strengths
Strengths	* Intuitive knowledge representation	* Data-driven learning	* Leverage strengths of both symbolic and
	* Reasoning capabilities	* Uncertainty handling	data-driven techniques
	* Scalability	* Flexibility	* Create more robust and flexible knowledge representation and reasoning
Data Reliance	Limited	High (data quality and quantity are crucial)	Can address limitations of both approaches
Uncertainty	Limited	Can quantify uncertainty	* Express confidence levels in inferences
			* Handle real-world data uncertainty
Learning	Limited	Enables learning from data	* Enrich nets with statistical data (e.g., probabilities on arcs)
			* Learn weights for relationships
Reasoning	Symbolic reasoning based on net structure	Statistical reasoning on net relationships	* Hybrid reasoning: symbolic + statistical
			for nuanced and robust reasoning
Challenges	Computational complexity for large nets	Computational complexity for large datasets	* Balancing symbolic and statistical aspects
			* Maintaining consistency between them



The screenshot shows a digital learning content slide. At the top left is the Parul University NAAC A++ logo. To the right is a header 'DIGITAL LEARNING CONTENT' with a circular icon containing various educational symbols. The main title is 'Probability and Bayes' theorem'. Below it, a list of bullet points explains the theorem and its application in AI. A section titled 'Understanding Probability:' follows, with three bullet points defining probability and its role in AI.

Probability and Bayes' theorem

- Bayes' theorem is also known as **Bayes' rule**, **Bayes' law**, or **Bayesian reasoning**, which determines the probability of an event with uncertain knowledge.
- In probability theory, it relates the conditional probability and marginal probabilities of two random events.
- Bayes' theorem was named after the British mathematician **Thomas Bayes**. The **Bayesian inference** is an application of Bayes' theorem, which is fundamental to Bayesian statistics.

Understanding Probability:

- Probability is the likelihood of an event occurring.
- It's expressed as a value between 0 (impossible) and 1 (certain).
- In AI, probability helps quantify the belief an AI system has about the truth of a statement or the likelihood of a particular outcome.

Bayes' theorem:

Bayes' theorem is also known as **Bayes' rule**, **Bayes' law**, or **Bayesian reasoning**, which determines the probability of an event with uncertain knowledge.

In probability theory, it relates the conditional probability and marginal probabilities of two random events.

Bayes' theorem was named after the British mathematician **Thomas Bayes**. The **Bayesian inference** is an application of Bayes' theorem, which is fundamental to Bayesian statistics.

It is a way to calculate the value of $P(B|A)$ with the knowledge of $P(A|B)$.

Example: If cancer corresponds to one's age then by using Bayes' theorem, we can determine the probability of cancer more accurately with the help of age.

Bayes' theorem can be derived using product rule and conditional probability of event A with known event B:

As from product rule we can write:

$$P(A \cap B) = P(A|B) P(B) \text{ or}$$

Similarly, the probability of event B with known event A:

$$P(A \cap B) = P(B|A) P(A)$$

Equating right hand side of both the equations, we will get:

$$P(A|B) = \frac{P(B|A) P(A)}{P(B)} \quad \dots(a)$$

The above equation (a) is called as **Bayes' rule** or **Bayes' theorem**. This equation is basic of most modern AI systems for **probabilistic inference**.

It shows the simple relationship between joint and conditional probabilities.

$P(A|B)$ is known as **posterior**, $P(B|A)$ is called the likelihood,

$P(A)$ is called the **prior probability**, probability of hypothesis before considering the evidence

$P(B)$ is called **marginal probability**, pure probability of an evidence.

DIGITAL LEARNING CONTENT	
	
Bayes' Theorem: Updating Beliefs with New Evidence	Applying Bayes' Theorem: A Medical Diagnosis
➤ Bayes' theorem is a powerful tool for revising beliefs based on new evidence. ➤ It allows AI systems to continuously learn and improve their decision-making. 	➤ Prior Probability: Doctor's initial belief about a patient's illness based on symptoms. ➤ Evidence: Test results (X-ray, blood work) ➤ Posterior Probability: Updated belief about the illness after considering test results.

Bayes' theorem is a powerful tool for **updating our beliefs** about the probability of an event (hypothesis) after we receive **new evidence**. It allows us to incorporate new information and refine our initial estimates.

Here's a breakdown of the key aspects:

- **Prior Probability ($P(A)$):** This represents our initial belief about the likelihood of an event (A) happening before considering any evidence. It's like our starting point.
- **New Evidence (Event B):** This is new information we receive that might influence our belief about event A.
- **Likelihood ($P(B|A)$):** This tells us how probable the new evidence (B) is, given that our initial belief (A) is true. In simpler terms, how likely is it to observe evidence B if

scenario A is actually true?

- **Posterior Probability ($P(A|B)$):** This is what we're solving for. It's the **updated probability** of event A occurring, after considering the new evidence (B). This is our refined belief after incorporating the new information.

The Formula:

$$P(A|B) = (P(B|A) * P(A)) / P(B)$$

- **$P(A|B)$:** Posterior probability (what we solve for)
- **$P(B|A)$:** Likelihood (probability of evidence B given A is true)
- **$P(A)$:** Prior probability (initial belief about A)
- **$P(B)$:** Total probability of evidence B (regardless of A)

Example:

Imagine you suspect your friend has a pet cat (A). You might have a prior belief based on general knowledge (e.g., $P(A) = 0.3$). Then, you hear faint meowing coming from their house (B). This is the new evidence.

- **Likelihood ($P(B|A)$):** How likely is it to hear meowing if your friend has a cat? This might be high (e.g., $P(B|A) = 0.8$).
- **Total probability of meowing ($P(B)$):** This considers all possibilities, not just cats. Maybe they have a neighbor with a cat or a recording (e.g., $P(B) = 0.1$).

Using the formula, you can calculate the updated probability (posterior probability) of your friend having a cat after hearing the meowing. This value will likely be higher than your initial belief (prior probability).

Applications:

Bayes' theorem has a wide range of applications across various fields:

- **Machine Learning:** It's a cornerstone of Bayesian learning algorithms that continuously adjust predictions based on new data.
- **Medical Diagnosis:** Doctors can use it to estimate the likelihood of a disease based on

symptoms and test results.

- **Spam Filtering:** Email filters leverage Bayes' theorem to classify emails as spam based on historical spam patterns.
- **Scientific Research:** Scientists can update their theories based on new experimental data.



DIGITAL LEARNING CONTENT



Applications of Probability and Bayes' Theorem in AI

- Machine Learning - Machine learning algorithms learn from data by identifying patterns and relationships.
- Spam Filtering - Spam filters use Bayes' theorem to categorize emails. They consider the presence of keywords and other indicators (evidence) associated with spam.
- Medical Diagnosis - AI systems can analyze patient data (symptoms, medical history) to estimate disease probabilities.
- Natural Language Processing (NLP) - NLP applications like chatbots or machine translation rely on probability.
- Robotics - Robots use probability to navigate their environment, avoid obstacles, and interact with objects.

1. Machine Learning:

- **Recommendation Systems:** These systems use probability to recommend products, movies, or music based on your past behavior and preferences. They can also use Bayes' theorem to update these recommendations as you interact with new content. For instance, if you consistently watch comedies on a streaming service, the system assigns a higher probability that you'll enjoy another comedy in the future.
- **Spam Filtering:** Email filters analyze incoming emails and calculate the probability of them being spam based on various factors like keywords and sender information. Bayes' theorem allows these filters to improve over time. If you consistently mark certain emails as spam, the filter updates its understanding of spam patterns and assigns a higher probability to similar emails being spam in the future.

2. Medical Diagnosis:

Doctors don't solely rely on symptoms to diagnose patients. They use probability to assess the likelihood of different diseases based on a patient's medical history and demographics. Bayes' theorem can be a powerful tool in this process.

- **Considering Prior Probability:** A doctor might consider the general prevalence of a disease in the population as the prior probability.

- **New Evidence:** Symptoms, test results, and family history act as new evidence (event B).
- **Likelihood:** The doctor would consider how likely these symptoms are for a particular disease ($P(B|A)$).
- **Refined Diagnosis:** Using Bayes' theorem, doctors can update their initial belief about the most likely disease (posterior probability) based on the new evidence.

3. Reliability Engineering:

- **Predictive Maintenance:** In various industries, probability and Bayes' theorem are used to predict equipment failures. They analyze historical maintenance data and equipment performance to calculate the probability of failure (event A).
- **New Evidence:** Sensor readings and inspection results can be considered as new evidence (event B).
- **Predicting Issues:** By calculating the likelihood of failure given these readings ($P(B|A)$), engineers can schedule maintenance before a critical breakdown occurs.

4. Finance and Insurance:

- **Credit Scoring:** Financial institutions use probability models to assess the creditworthiness of loan applicants. These models consider income, credit history, and other factors to determine the probability of loan repayment (event A).
- **Fraud Detection:** Banks use transaction data and customer behavior to calculate the probability of fraudulent activity (event B). Bayes' theorem allows them to update this probability in real-time as new transactions occur.

5. Robotics and Self-Driving Cars:

- **Navigation and Obstacle Avoidance:** Robots and self-driving cars use probability to navigate their environment. They create maps with probabilities assigned to different locations representing free space, obstacles, or potential hazards. Sensor data (like LiDAR) acts as new evidence (event B). Bayes' theorem allows them to update the probability of encountering obstacles in real-time (posterior probability), leading to safer and more efficient navigation.